

NUMEN MVP MASTER ROADMAP

Purpose

This document is the execution plan for the next build sequence of Numen/Nutrition Autopilot. It is written for direct use by Claude Code inside the repo.

This is **not** a greenfield build. This is a structured continuation of an existing working system. Build incrementally. Preserve the current pipeline. Prioritize scientific rigor and operational usability.

Operating Mode

Full-Run Mode

When asked to "run full," execute all sprints in sequence **without re-planning the whole system each time.**

For each sprint:

1. Inspect current extension points only (schema/API/UI/tests impacted by that sprint)
2. Implement smallest safe version of required scope
3. Add/extend tests before UI polish
4. Run validations + regressions
5. Update handoff docs
6. Gate pass/fail
7. Continue to next sprint only if gate passes

If a gate fails:

- stop advancing to the next sprint
- stabilize current sprint
- document blocker and exact next steps

Build Style

- Engineering-first
- Incremental
- No architecture recap unless needed for a decision

- No broad rewrites of stable modules
 - No fake certainty
 - No hidden state changes
-

Global Non-Negotiables

Scientific Rigor

- Preserve deterministic nutrient math behavior
- Preserve provenance tracking for nutrient values
- Preserve immutable label freeze semantics
- Preserve auditability (traceable inputs, weights, lots, nutrient sources)
- No silent substitutions or silent assumptions
- Document assumptions in `ASSUMPTIONS.md`

Pipeline Safety

Do not break:

- SOT import
- Instacart import
- enrichment pipeline
- schedule flow
- Fed/Skip serving flow
- immutable label freeze
- label generation

MVP Scope Discipline

This roadmap is for:

- 1 client
- 1 small kitchen
- PWA-first chef/admin usage

Do not build:

- multi-kitchen orchestration
- complex roles/permissions
- native mobile priority work (Expo remains secondary)
- speculative autonomous planner behavior that bypasses review

Required Handoff Files (update every sprint)

- PROGRESS.md
- OPEN_ISSUES.md
- SCIENTIFIC_RISKS.md
- MORNING_HANDOFF.md

Each sprint must update these with:

- shipped work
- remaining work
- test status
- changed files summary
- migration notes (if any)
- exact commands to run locally
- next sprint recommendation

Validation Gate (must pass before advancing)

A sprint is considered gated/passed only if all are true:

- tests green (including new sprint tests)
- typecheck green
- lint green
- web build green (if UI touched)
- no regression in import → enrichment → schedule → label freeze
- SCIENTIFIC_RISKS.md updated with any new risk/assumption
- migration/apply status documented if schema changed

If any fail, fix or clearly document blocker before continuing.

Rollback / Safety Rules

- Work on a feature branch unless explicitly told otherwise
- Commit in logical checkpoints
- Do not squash away scientific test history during active sprinting
- If a migration is risky, prefer additive schema changes over destructive changes

- Never mutate historical label snapshots in debug/recompute tools
-

Sprint Sequence (Dependency-Ordered)

Sprint 1 — Inventory Intelligence + Instacart Mapping UX + Substitution Engine (MVP)

Objective

Improve operational reliability and speed in kitchen/admin workflows by adding:

1. inventory intelligence (projections / allocation / reorder signals)
2. fast Instacart mapping review UX
3. practical substitution engine with scientific rigor + auditability

Scope (Required)

A. Inventory Intelligence

Build:

- projected availability by date/time window using:
 - inventory lots
 - meal schedules
 - batch prep demand
 - allocations/reservations
- 7-day inventory demand forecast:
 - by ingredient/component family
 - required vs available vs shortfall
- par levels / reorder thresholds (MVP-simple, configurable)
- allocation visibility:
 - on-hand / allocated / available / projected balance
- waste analytics (MVP-light):
 - quantities + counts by reason and ingredient/component

UX:

- surface high-signal states:
 - shortage this week
 - critical today
 - overallocated
 - expiring soon
- use clear badges/status colors (do not bury risk in tables)

B. Instacart Mapping UX

Build:

- unmapped line item queue (sortable/filterable)
- candidate match suggestions ranked by confidence:
 - UPC exact
 - normalized name similarity
 - brand match
 - size/unit similarity
 - historical accepted mappings
- one-click actions:
 - approve suggested match
 - search/select different existing SKU/ingredient
 - create new SKU + link to canonical ingredient
 - mark pantry/non-tracked with reason
- mapping memory / learned mappings
- dry-run + apply workflow

Requirements:

- preserve import idempotency / duplicate protections
- preserve provenance/confidence metadata
- record mapping resolution source (manual vs auto-suggest accepted)

C. Substitution Engine (MVP)

Build substitution suggestions for planned meals/batches using:

- same component family
- allergen/exclusion compatibility
- inventory availability
- prep readiness / batch status
- nutrient delta score
- optional flavor compatibility tags (if available)

Add:

- explainable ranking ("why suggested")
- preview before apply (before/after nutrient delta)
- apply action for future planned meals/batches only
- audit event/reason logging

Constraints:

- no silent substitutions
- no changes to served/frozen labels
- warnings when uncertainty is high

D. Scientific QA + Regression Hardening

Add tests for:

- inventory allocation invariants
- projection logic correctness
- mapping memory + idempotency preservation
- substitution constraints/ranking/deltas
- deterministic nutrient recalculation post-substitution
- rounding + calorie sanity after substitution
- no impact on served/frozen labels

Extend QA/admin exceptions to show:

- unresolved mapping queue counts
- overallocated inventory
- substitution warnings
- projected shortages affecting scheduled meals

Sprint 1 Acceptance Criteria

- Instacart mapping review queue is usable end-to-end
 - inventory page surfaces projected shortages and overallocation
 - substitutions can be previewed and applied to future plans only
 - scientific and regression tests pass
 - no label freeze regressions
-

Sprint 2 — Batch Yield Calibration + Cook/Chill QC Telemetry (MVP)

Objective

Replace static yield assumptions with measured production data and capture QC checkpoints that improve scientific rigor and planning accuracy.

Scope (Required)

A. Batch Yield Calibration

Build:

- yield profile history (expected %, actual %, variance, sample count, confidence)
- calibration dimensions (MVP):
 - component/batch type
 - method
 - optional cut/form factor
- suggested yield update workflow (human accept required)
- outlier handling (flag + exclude from calibration proposal)

B. Cook/Chill QC Telemetry

Build:

- checkpoint logging:
 - cook start/end
 - target vs actual temp
 - chill start/end
 - chill compliance
 - hold status/notes
- step requirements before state transitions (MVP-simple)
- issue flags:
 - temp miss
 - chill time exceeded
 - missing/late checkpoint
 - manual override with reason

C. Batch Variance Analytics

Build:

- expected vs actual yield summaries
- trend over time
- repeated variance flags
- top problematic batches/components
- light contributor views (method / notes / patterns if available)

D. Planning Integration

Planning/projection logic should:

- prefer approved calibrated yields
- fall back to defaults if confidence insufficient
- explicitly show basis used (default vs calibrated)

No silent replacement of assumptions.

Scientific QA + Tests

Add tests for:

- calibration proposal logic
- outlier rejection
- calibrated/default selection
- checkpoint gating invariants
- QC override requirements
- deterministic planning math with calibrated yields
- no impact on frozen labels
- calorie/rounding sanity remains green

Extend QA/admin exceptions to show:

- missing required checkpoints
- repeated yield variance issues
- unreviewed calibration proposals
- QC overrides requiring review

Sprint 2 Acceptance Criteria

- checkpoint data can be logged in kitchen flow
- calibrated yields are proposed and reviewable
- planning can use calibrated yields with explicit traceability
- QA surfaces QC misses and calibration risks

- tests/regressions green
-

Sprint 3 — Menu Composer + Weekly Prep Optimizer + Sauce Matrix (MVP)

Objective

Operationalize composition-first planning:

- standardized components (protein/base/veg/sauce)
- minimized prep complexity
- personalized flavor via sauces

Scope (Required)

A. Menu Composer

Build a composition-first planner:

- choose protein/base/veg/sauce
- set gram targets / presets
- macro/calorie preview
- save reusable composition template
- compatibility guidance (tags)
- allergen/incompatibility warnings
- flavor swaps via sauce variants without duplicating core prep assumptions

B. Weekly Prep Optimizer (MVP heuristic)

Build:

- 7-day component demand rollup
- cooked quantity required
- raw quantity estimate via approved/default yields
- bundling recommendations (share neutral batches across multiple meals)
- prep draft generation (review before commit)
- shortage/blocker surfacing integrated with inventory intelligence

C. Sauce Matrix

Build:

- sauce matrix UI:
 - flavor family
 - macro variant
 - compatibility tags
- default sauce assignments for component combinations/templates
- portion presets (5g/10g/15g/20g...) with macro delta preview
- flavor rotation support (MVP-light) to reduce repetition

D. Schedule Integration (coexist recipe-first + composition-first)

Build compatibility path so planned meals can use composition templates without breaking existing recipe-based schedule flow.

Requirements:

- preserve label freeze and nutrient traceability
- preserve QA visibility
- if compatibility layer is needed, implement and document clearly

Scientific QA + Tests

Add tests for:

- composition macro aggregation determinism
- sauce portion/variant effects
- yield-adjusted prep rollup calculations
- inventory shortage detection in planning
- compatibility/allergen warning logic
- recipe-first + composition-first coexistence behavior
- no regression in label freeze pipeline

Extend QA/admin warnings for:

- composition templates with estimated nutrient inputs
- rollups using low-confidence yields
- shortages blocking prep draft generation

Sprint 3 Acceptance Criteria

- composition templates can be created and previewed
- weekly prep rollup produces usable batch suggestions
- sauce matrix supports practical personalization

- existing schedule/label flow still works
 - tests/regressions green
-

Sprint 4 — Client Biometrics Timeline + DEXA/Bloodwork/CGM Ingestion Layer (MVP)

Objective

Create a structured client data layer for time-series biometrics and document ingestion with auditability and future parsing hooks.

This sprint is storage/indexing/timeline-first. Do not fake parsers.

Scope (Required)

A. Client Biometrics Timeline

Build:

- biometric snapshots (date-stamped)
 - height
 - weight
 - body fat % (optional)
 - lean mass (optional)
 - notes/source
- timeline UI:
 - chronological records
 - latest values summary
 - add/edit snapshots
 - missing/irregular data indicators
- trend preview (MVP-light):
 - up/down/stable indicators

B. Document Ingestion Layer

Build structured client document records:

- type (`DEXA`, `BLOODWORK`, `CGM`, `OTHER`)
- date collected
- date uploaded
- source/provider (optional)
- tags
- file linkage/storage metadata
- parsing status (`not_started`, `queued`, `parsed_partial`, `verified`, `failed`)

Add upload/management UI:

- upload file
- assign type/date
- view metadata
- mark verified
- add notes
- filter/search by type/date/status

C. Parsed Metrics Scaffold (future-ready, MVP-safe)

Create normalized metric series storage:

- metric key
- value + unit
- observed_at
- source document reference
- confidence/verification status

Add manual entry path for common metrics (MVP):

- fasting glucose
- LDL/HDL/triglycerides
- HbA1c
- body fat % / lean mass
- optional resting HR

No fake parsing. Keep explicit manual/placeholder workflow when parser is absent.

D. QA/Admin Integration

Surface:

- missing recent biometrics
- unverified documents
- failed parsing attempts (if any)
- stale metric data affecting planning confidence (MVP rule-based)

Scientific QA + Tests

Add tests for:

- time-series ordering and latest selection
- document metadata integrity
- metric linkage to source documents
- parsing status transitions
- verification flags
- no impact on existing meal/label pipeline

No silent unit assumptions; keep source + verification states explicit.

Sprint 4 Acceptance Criteria

- client timeline usable
 - document upload/metadata workflow usable
 - metrics can be stored with provenance/verification
 - QA warns on stale/unverified data
 - tests/regressions green
-

Sprint 5 — Audit Trace + Reproducibility + Ops Control Tower (MVP)

Objective

Make the system operationally and scientifically legible:

1. audit trace visibility
2. reproducibility/debug tooling
3. high-signal admin control tower

Scope (Required)

A. Label/Meal Audit Trace Viewer

For a selected served meal/label snapshot, show:

- schedule/service event context
- freeze-time recipe/composition inputs
- ingredients/components + gram weights
- inventory lots consumed
- nutrient provenance for contributing values
- calculation summary (human-readable)
- QA warnings/confidence flags present at freeze time

Requirements:

- preserve immutable snapshot semantics
- distinguish stored snapshot vs current data if both are shown
- no historical mutation

B. Reproducibility / Debug Tools (non-destructive)

Build:

- recompute preview diff (historical snapshot vs current computed result)
- delta explanation (mapping changes, nutrient source changes, yields, etc.)
- freeze-time integrity checks (references present and internally consistent)

No mutation of snapshots. Diff-only.

C. Ops Control Tower Dashboard

Build a single high-signal dashboard with sections:

- Today:
 - meals due
 - batches due/active/blocked
 - shortages/overallocations
 - expiring inventory
- Scientific QA:
 - unresolved verification tasks
 - estimated/inferred nutrient counts
 - missing provenance/incomplete coverage
 - substitution warnings pending review
 - calibration/QC override review items
- Client Data Readiness:
 - stale biometrics/docs
 - unverified uploads
- System Reliability (MVP-light):
 - failed imports
 - failed enrichment jobs

- stuck statuses
- Attention Queue:
 - top actionable issues with direct links

D. Exportable Review Packs (MVP)

Print/export-friendly views:

- Daily Ops Summary
- Scientific QA Summary
- Meal Audit Summary (single meal/label)

Browser print is acceptable. Prioritize correctness and readability.

Scientific QA + Tests

Add tests for:

- audit trace integrity and provenance display logic
- recompute preview diff behavior (non-destructive)
- dashboard aggregation correctness
- attention queue deterministic prioritization (if scoring used)
- no mutation of label snapshots during debug paths
- regression coverage for current pipeline

Sprint 5 Acceptance Criteria

- audit trace viewer explains freeze-time label generation
- recompute diff is safe and useful
- control tower surfaces highest-risk issues clearly
- export/print review views are usable
- tests/regressions green

Full-Run Execution Protocol (Claude Code)

Start Sequence

1. Read this file.
2. Detect current codebase state and which sprint is next.
3. If no sprint-tracking marker exists, start at Sprint 1.

4. Create/update a sprint tracking file:
 - `ROADMAP_STATUS.md`
 - mark current sprint, status, gate result, notes
5. Execute current sprint.
6. Run gate.
7. If gate passes, advance `ROADMAP_STATUS.md` and continue.
8. If gate fails, stop advancement and document blockers precisely.

Sprint Tracking Format (use this exact shape)

Maintain `ROADMAP_STATUS.md` with:

- Current sprint
- Status (`not_started`, `in_progress`, `gated_pass`, `blocked`)
- Last completed gate timestamp
- Blockers
- Next action

Implementation Discipline

For each sprint:

- inspect extension points only
- implement data model/API changes first (where needed)
- add tests for invariants before UI polish
- implement UI
- run validations/regressions
- update docs
- commit checkpoint(s)

Testing / Validation Commands

Discover and use repo-native commands from `package.json` / workspace scripts. Document exact commands in `MORNING_HANDOFF.md`.

Do not invent command names. Verify before writing docs.

Stop/Ship Criteria for MVP Funding Demo

The roadmap can be considered demo-ready when all are true:

- Sprints 1–5 gated pass
 - core pipeline remains stable
 - chef/admin PWA workflows are usable in real kitchen simulation
 - scientific QA and audit trace are visible and trustworthy
 - control tower surfaces operational and scientific risk clearly
 - handoff docs are current and credible
-

Default Prioritization Rule (when time/complexity tradeoffs appear)

Choose in this order:

1. Scientific correctness / auditability
2. Pipeline safety
3. Chef/admin operational usability
4. Speed/automation convenience
5. UI polish

If a tradeoff is made, document it in [SCIENTIFIC_RISKS.md](#) or [OPEN_ISSUES.md](#).

Final Instruction to Claude

Build directly against current context and current repo state. Do not reset. Do not genericize. Ship the highest-leverage increment that preserves scientific rigor and operational momentum.